

Sistemas Operativos (75.08): Lab Kernel

Lavandeira, Lucas (#98042)
lucaslavandeira@gmail.com

Rozanec, Matias (#97404)
rozanecm@gmail.com

29.jun.2018

Parte I

Código desalojo

Listing 1: contador.c

```
1  #include "decls.h"
2
3  #define COUNTLEN 20
4  #define TICKS (1ULL << 15)
5  #define DELAY(x) (TICKS << (x))
6  #define USTACK_SIZE 4096
7
8  static volatile char *const VGABUF = (volatile void *) 0xb8000;
9
10 static uintptr_t esp;
11 static uint8_t stack1[USTACK_SIZE] __attribute__((aligned(4096)));
12 static uint8_t stack2[USTACK_SIZE] __attribute__((aligned(4096)));
13
14 static void exit() {
15     uintptr_t tmp = esp;
16     esp = 0;
17     if (tmp)
18         task_swap(&tmp);
19 }
20
21 static void yield() {
22     if (esp)
23         task_swap(&esp);
24 }
25
26 static void contador_yield(unsigned lim, uint8_t linea, char color) {
27     char counter[COUNTLEN] = {'0'}; // ASCII digit counter (RTL).
28
29     while (lim-->0) {
30         char *c = &counter[COUNTLEN];
31         volatile char *buf = VGABUF + 160 * linea + 2 * (80 - COUNTLEN);
32
33         unsigned p = 0;
34         unsigned long long i = 0;
35
36         while (i++ < DELAY(6)) // Usar un entero menor si va demasiado lento.
37             ;
38
39         while (counter[p] == '9') {
40             counter[p++] = '0';
41         }
42
43         if (!counter[p]) {
44             counter[p] = '1';
45         }
46
47         while (c-- > counter) {
48             *buf++ = *c;
49             *buf++ = color;
50         }
51
52         yield();
53     }
54     // exit();
55 }
56
57 void contador_run() {
58     // Inicializar al *tope* de cada pila.
59     uintptr_t *a = (uintptr_t*) (stack1 + USTACK_SIZE);
60     uintptr_t *b = (uintptr_t*) (stack2 + USTACK_SIZE);
61
62     *--a = 0x2F;
63     *--a = 0;
64     *--a = 100;
65
66     *--b = 0x4F;
67     *--b = 1;
68     *--b = 10;
69     *--b = (uintptr_t) exit;
70     *--b = (uintptr_t) contador_yield;
71     *--b = 0;
72     *--b = 0;
73     *--b = 0;
74     *--b = 0;
75
76     esp = (uintptr_t) b;
77
78     task_exec((uintptr_t) contador_yield, (uintptr_t) a);
79 }
80
81
82
83 static void contador1() {
```

```

84     contador_yield(50000000, 2, 0x2F);
85 }
86
87 static void contador2() {
88     contador_yield(50000000, 3, 0x6F);
89 }
90
91 static void contador3() {
92     contador_yield(50000000, 4, 0x4F);
93 }
94
95 void contador_spawn() {
96     spawn(contador1);
97     spawn(contador2);
98     spawn(contador3);
99 }

```

Listing 2: kern0.c

```

1  #include "decls.h"
2  #include "multiboot.h"
3  #include "string.h"
4  #include "interrupts.h"
5  #include "sched.h"
6
7  #define USTACK_SIZE 4096
8
9  static uint8_t stack1[USTACK_SIZE] __attribute__((aligned(4096)));
10 static uint8_t stack2[USTACK_SIZE] __attribute__((aligned(4096)));
11
12 void two_stacks_c() {
13     // Inicializar al *tope* de cada pila.
14     uintptr_t *a = (uintptr_t*) (stack1 + USTACK_SIZE);
15     uintptr_t *b = (uintptr_t*) (stack2 + USTACK_SIZE);
16
17     // Preparar, en stack1, la llamada:
18     //   vga_write("vga_write() from stack1", 15, 0x57);
19
20     *--a = 0x57;
21     *--a = 15;
22     *--a = (uintptr_t) "vga_write() from stack1";
23
24     // AYUDA 1: se puede usar alguna forma de pre- o post-
25     // incremento/decremento, según corresponda:
26     //
27     //     *(a++) = ...
28     //     *(++a) = ...
29     //     *(a--) = ...
30     //     *--a = ...
31
32     // AYUDA 2: para apuntar a la cadena con el mensaje,
33     // es suficiente con el siguiente cast:
34     //
35     //     ... a ... = (uintptr_t) "vga_write() from stack1";
36
37     // Preparar, en s2, la llamada:
38     // vga_write("vga_write() from stack2", 16, 0xD0);
39
40     // AYUDA 3: para esta segunda llamada, usar esta forma de
41     // asignación alternativa:
42     b -= 3;
43     b[0] = (uintptr_t) "vga_write() from stack2";
44     b[1] = 16;
45     b[2] = 0xD0;
46
47     // Primera llamada usando task_exec().
48     task_exec((uintptr_t) vga_write, (uintptr_t) a);
49
50     // Segunda llamada con ASM directo. Importante: no
51     // olvidar restaurar el valor de %esp al terminar, y
52     // compilar con: -fasm -fno-omit-frame-pointer.
53     asm("movl %0, %%esp; call *%1; movl %%ebp, %%esp"
54         : /* no outputs */
55         : "r"(b), "r"(vga_write));
56 }
57
58 void kmain(const multiboot_info_t *mbi) {
59     int8_t linea;
60     uint8_t color;
61
62     vga_write("kern2 loading.....", 8, 0x70);
63
64     if (mbi->flags & MULTIBOOT_INFO_CMDLINE) {
65         char buf[256] = "cmdline: ";
66         char *cmdline = (void *) mbi->cmdline;
67
68         strlcat(buf, cmdline, 256);
69         vga_write(buf, 9, 0x07);
70     }
71
72     char mem[256] = "Physical memory: ";

```

```

73     char tmp[64] = "";
74
75     if (fmt_int(mbi->mem_upper - mbi->mem_lower, tmp, sizeof tmp)) {
76         strlcat(mem, tmp, sizeof mem);
77         strlcat(mem, "MiB total", sizeof mem);
78     }
79
80     vga_write(mem, 10, 0x07);
81
82     /* A remplazar por una llamada a two_stacks(),
83      * definida en stacks.S.
84      */
85     /*
86     vga_write("vga_write() from stack1", 12, 0x17);
87     vga_write("vga_write() from stack2", 13, 0x90);
88     */
89     two_stacks();
90     two_stacks_c();
91
92     contador_run();
93
94     vga_write("antes del 2", 18, 0xE0);
95     vga_write2("Funciona vga_write2?", 18, 0xE0);
96
97     /* From Ej: kern2-task:
98      * Por último, como "bootstrap" del planificador, se necesita una
99      * llamada a la función sched_init() desde kmain(), antes de las llamadas
100      * a idt_init()/irq_init(). Esto se necesita para que haya una
101      * tarea inicial en ejecución.
102      */
103     sched_init();
104     // Código ejercicio kern2-idt.
105     idt_init(); // (a)
106     irq_init();
107     asm("int3"); // (b)
108
109     asm("div %4"
110         : "=a"(linea), "=c"(color)
111         : "0"(18), "1"(0xE0), "b"(1), "d"(0));
112
113     vga_write("Funciona vga_write2?", linea, color);
114
115     contador_spawn();
116     asm("hlt");
117 }

```

Listing 3: idt_entry.S

```

1  #define PIC1 0x20
2  #define ACK_IRQ 0x20
3
4
5  .globl breakpoint
6  breakpoint:
7      // (1) Guardar registros.
8      push %eax
9      push %edx
10     push %ecx
11
12
13     // (2) Preparar argumentos de la llamada.
14     mov $breakpoint_msg, %eax
15     mov $14, %edx
16     mov $0xE0, %ecx
17
18     // (3) Invocar a vga_write2()
19     call vga_write2
20
21     // (4) Restaurar registros.
22     pop %ecx
23     pop %edx
24     pop %eax
25
26     // (5) Finalizar ejecución del manejador.
27     iret
28
29
30 .globl ack_irq
31 ack_irq:
32     // Indicar que se manejó la interrupción.
33     movl $ACK_IRQ, %eax
34     outb %al, $PIC1
35     iret
36
37
38 .globl timer_asm
39 timer_asm:
40     // Guardar registros e invocar handler
41     pusha
42     call timer
43

```

```

44 // Ack *antes* de llamar a sched()
45 movl $ACK_IRQ, %eax
46 outb %al, $PIC1
47
48 // Llamada a sched con argumento
49 push %esp
50 call sched
51
52 // Retornar (si se volvi6 de sched)
53 addl $4, %esp
54 popa
55 iret
56
57
58 .globl divzero
59 divzero:
60 // (1) Guardar registros.
61 push %eax
62 push %edx
63 push %ecx
64
65 movl $divzero_msg, %eax
66 movl $17, %ecx
67 call vga_write_cyan
68
69 // (4) Restaurar registros.
70 pop %ecx
71 pop %edx
72 pop %eax
73 jmp ack_irq
74
75 .globl keyboard_asm
76 keyboard_asm:
77 // (1) Guardar registros.
78 push %eax
79 push %edx
80 push %ecx
81
82 call keyboard
83
84 // (4) Restaurar registros.
85 pop %ecx
86 pop %edx
87 pop %eax
88 jmp ack_irq
89
90 .data
91 breakpoint_msg:
92 .asciz "Hello, breakpoint"
93
94
95 divzero_msg:
96 .asciz "Se divide por ++ebx"

```

Listing 4: sched.c

```

1 #include "decls.h"
2 #include "sched.h"
3
4 #define MAX_TASK 10
5
6 static struct Task Tasks[MAX_TASK];
7 static struct Task *current;
8
9 void sched_init() {
10     current = &Tasks[0];
11
12     for (int i = 0; i < MAX_TASK; i++) {
13         Tasks[i].status = FREE;
14         Tasks[i].frame = 0;
15     }
16
17     current->status = RUNNING;
18 }
19
20 void spawn(void (*entry)(void)) {
21     for(int i = 0; i < MAX_TASK; i++) {
22         if (Tasks[i].status == FREE) {
23             Tasks[i].status = READY;
24
25             size_t frame_size = sizeof(struct TaskFrame);
26             uint8_t* stack_top = &Tasks[i].stack[4096];
27
28             Tasks[i].frame = stack_top - frame_size;
29
30             Tasks[i].frame->edi = 0;
31             Tasks[i].frame->esi = 0;
32             Tasks[i].frame->ebp = 0;
33             Tasks[i].frame->esp = 0;
34             Tasks[i].frame->eax = 0;
35             Tasks[i].frame->ecx = 0;

```

```

36         Tasks[i].frame->edx = 0;
37         Tasks[i].frame->ebx = 0;
38
39         Tasks[i].frame->eflags = 0x0200; // flag IF = 1
40
41         Tasks[i].frame->cs = 8; // Multiboot siempre pone '8'
42                                // como CS (ver interrupts.c)
43         Tasks[i].frame->eip = entry;
44         return;
45     }
46 }
47 }
48
49 void sched(struct TaskFrame *tf) {
50     struct Task *new = 0;
51     struct Task *old = current;
52
53     int running_pos = 0;
54     for (int i = 0; i < MAX_TASK; i++) {
55         if (Tasks[i].status == RUNNING) {
56             running_pos = i;
57             break;
58         }
59     }
60
61     int pos = running_pos;
62     while (!new) {
63         if (Tasks[pos].status == READY) {
64             new = &Tasks[pos];
65         }
66         pos++;
67
68         if (pos == MAX_TASK) {
69             pos = 0;
70         }
71     }
72
73     old->status = READY;
74     old->frame = tf;
75
76     new->status = RUNNING;
77     current = new;
78     asm("movl %0, %%esp\n"
79         "popa\n"
80         "iret\n"
81         :
82         : "g"(current->frame)
83         : "memory");
84 }

```

Listing 5: sched.h

```

1 void sched_init();
2 void contador_spawn();
3
4
5 enum TaskStatus {
6     FREE = 0,
7     READY,
8     RUNNING,
9     DYING,
10 };
11
12 struct TaskFrame {
13     uint32_t edi;
14     uint32_t esi;
15     uint32_t ebp;
16     uint32_t esp;
17     uint32_t ebx;
18     uint32_t edx;
19     uint32_t ecx;
20     uint32_t eax;
21     /* below here defined by x86 hardware */
22     uint32_t eip;
23     uint16_t cs;
24     uint16_t padding;
25     uint32_t eflags;
26 } __attribute__((packed));
27
28
29 struct Task {
30     uint8_t stack[4096];
31     enum TaskStatus status;
32     struct TaskFrame *frame;
33 };

```